




**Samueli**  
Computer Science



# CS32: Introduction to Computer Science II

## **Discussion Week 8**

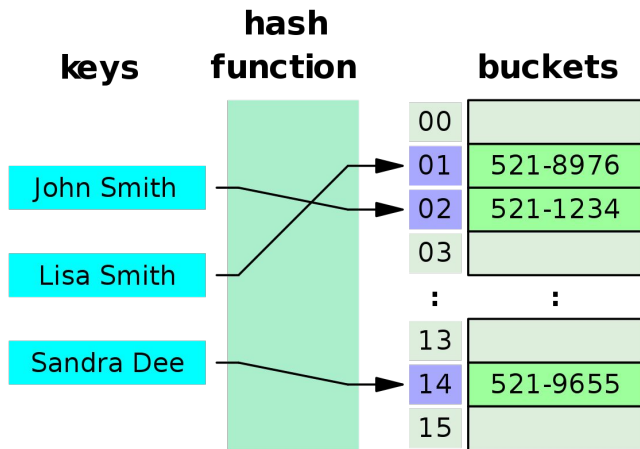
Yichao (Joey)

Feb. 28, 2019

- Hash Table

- Homework 4 is due 11:00 PM Tuesday, March 3.

# Hash Tables



In computing, a **hash table** (hash map) is a data structure that implements an **associative array abstract data type**, a structure that can **map keys to values**.

A hash table uses a **hash function** to compute an *index*, also called a *hash code*, into an array of *buckets* or *slots*, from which the desired value can be found.

The idea of hashing is to distribute the entries (key/value pairs) across an array of *buckets*. Given a key, the algorithm computes an *index* that suggests where the entry can be found:

```
index = f(key, array_size)
```

Often this is done in two steps:

```
hash = hashfunc(key)  
index = hash % array_size
```

In this method, the *hash* is independent of the array size, and it is then *reduced* to an index (a number between 0 and `array_size - 1`) using the **modulo operator** (%).

# Hash Tables

## Open vs Closed Addressing

### Open Addressing

Also known as **closed hashing**.

Collisions are dealt with by searching for **another empty buckets** within the hash table array itself.

Benefits:

- Better memory locality and cache performance. All elements laid out linearly in memory.
- Performs better than closed addressing when the number of keys is known in advance and the churn is low.

0 :  
1 : ①  
2 : ②  
3 : ②  
4 : ②  
5 : ④  
6 :  
7 : ⑦  
8 : ⑦  
9 : ⑨

### Closed Addressing

Also known as **open hashing**.

A key is always stored in the bucket it's hashed to. Collisions are dealt with using **separate data structures** on a per-bucket basis.

0 :  
1 : ①  
2 : ②②②  
3 :  
4 : ④  
5 :  
6 :  
7 : ⑦⑦  
8 :  
9 : ⑨

Benefits:

- Easier removal (no need for deleted markings)
- Typically performs better with high load factor.

- Insert
- Remove
- Search
  
- The complexity depends on your hash tables.
- Closed Hashing
  - Fixed number of buckets
  - All operations are  $O(n)$  with a small constant of proportionality
- Open Hashing
  - Consider  $\#entries / \#buckets$
  - Almost  $O(1)$  for all operations

Pretty much any time you want to map keys to values with constant time for lookup/add/remove. If you find yourself **looping through a list to find an element**, think if there's a way to store the elements with keys in a **hash table** (aka hash map, aka dictionary) instead.

```
1. class Contacts {
2.     List<Person> contacts = new ArrayList<Person>();
3.
4.     void addContact(Person p) {
5.         contacts.add(p);
6.     }
7.     // O(n) time
8.     Person getContact(String phoneNumber) {
9.         for (Person p : contacts) {
10.            if (p.getPhoneNumber().equals(phoneNumber))
11.                return p;
12.        }
13.        return null;
14.    }
15.    ...
16. }
```

```
1. class Contacts {
2.     List<Person> contacts = new ArrayList<Person>();
3.     Map<String, Person> phoneNumberToPerson = new
4.         HashMap<String, Person>();
5.
6.     void addContact(Person p) {
7.         contacts.add(p); // assuming this list is still
8.         needed somewhere else
9.         phoneNumberToPerson.put(p.getPhoneNumber(), p);
10.    }
11.    // O(1) time!
12.    Person getContact(String phoneNumber) {
13.        return phoneNumberToPerson.get(phoneNumber);
14.    }
15.    ...
16. }
```



**Que – 1.** Given the following input (4322, 1334, 1471, 9679, 1989, 6171, 6173, 4199) and the hash function  $x \bmod 10$ , which of the following statements are true?

- i. 9679, 1989, 4199 hash to the same value
- ii. 1471, 6171 has to the same value
- iii. All elements hash to the same value
- iv. Each element hashes to a different value

- (A) i only
- (B) ii only
- (C) i and ii only
- (D) iii or iv

# Hash Tables

## Simple question

**Que – 2.** The keys 12, 18, 13, 2, 3, 23, 5 and 15 are inserted into an initially empty hash table of length 10 using open addressing with hash function  $h(k) = k \bmod 10$  and linear probing. What is the resultant hash table?

|   |    |
|---|----|
| 0 |    |
| 1 |    |
| 2 | 2  |
| 3 | 23 |
| 4 |    |
| 5 | 15 |
| 6 |    |
| 7 |    |
| 8 | 18 |
| 9 |    |

(A)

|   |    |
|---|----|
| 0 |    |
| 1 |    |
| 2 | 12 |
| 3 | 13 |
| 4 |    |
| 5 | 5  |
| 6 |    |
| 7 |    |
| 8 | 18 |
| 9 |    |

(B)

|   |    |
|---|----|
| 0 |    |
| 1 |    |
| 2 | 12 |
| 3 | 13 |
| 4 | 2  |
| 5 | 3  |
| 6 | 23 |
| 7 | 5  |
| 8 | 18 |
| 9 | 15 |

(C)

|   |           |
|---|-----------|
| 0 |           |
| 1 |           |
| 2 | 12, 2     |
| 3 | 13, 3, 23 |
| 4 |           |
| 5 | 5, 15     |
| 6 |           |
| 7 |           |
| 8 | 18        |
| 9 |           |

(D)

# Hash Tables

## Simple question

**Que – 3.** A hash table of length 10 uses open addressing with hash function  $h(k)=k \bmod 10$ , and linear probing. After inserting 6 values into an empty hash table, the table is as shown below.

Which one of the following choices gives a possible order in which the key values could have been inserted in the table?

- (A) 46, 42, 34, 52, 23, 33
- (B) 34, 42, 23, 52, 33, 46
- (C) 46, 34, 42, 23, 52, 33
- (D) 42, 46, 33, 23, 34, 52

|   |    |
|---|----|
| 0 |    |
| 1 |    |
| 2 | 42 |
| 3 | 23 |
| 4 | 34 |
| 5 | 52 |
| 6 | 46 |
| 7 | 33 |
| 8 |    |
| 9 |    |

# Hash Tables

## FNV-1

- Hash functions: Take a “key” and map it to a number
- Requirement for hash function: should return the same value for the same key
- Good hash functions:
  - Spreads out the values: two different key are likely to results in different hash values. → **Avoid confliction**
  - Compute each value quickly.
- Example: **FNV-1**



```
unsigned int FNV-1(string s) {  
    unsigned int h = 2166136261U; ← offset_basis  
    for (int k = 0; k != s.size(); k++)  
    {  
        h += s[k];  
        h *= 16777619; ← FNV_prime  
    }  
    return h;  
}
```


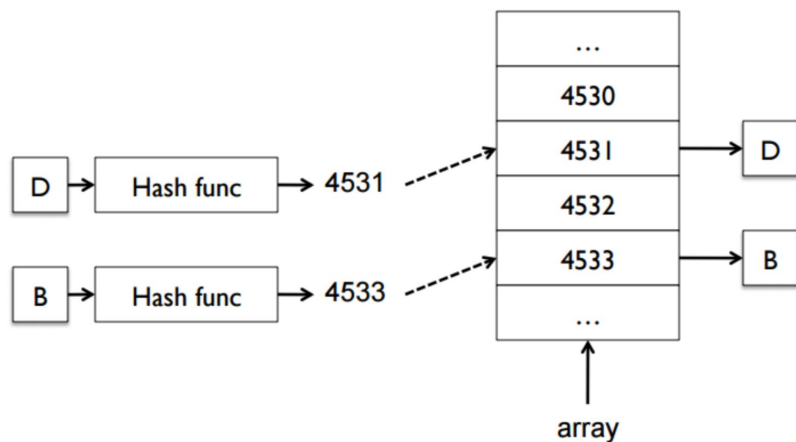
Fowler–Noll–Vo (FNV) is a non-cryptographic hash function created by Glenn Fowler, Landon Curt Noll, and Kiem-Phong Vo.  
<http://www.isthe.com/chongo/tech/comp/fnv/#FNV-param>

# Hash Tables

## Examples

- Example: Use a hash table to store people.
- Use a linked list to collision in the hash function.

*"You should almost **NEVER** assume that collisions are impossible!!!" --David Smallberg*



# Hash Tables Problem

## Word counting and sorting in a document

---

- Question: We have  $n$  words in a document, whose vocabulary size is  $v$ . Count top  $k$  frequent words in a document.

Two steps: counting + sorting

- The most efficient way to count the frequency for all words takes  $O(\underline{n})$  time complexity.
- After getting the frequency of each word, the most efficient way to get the top  $k$  frequent words takes  $O(\underline{v \log k})$  time complexity.
- Totally the entire procedure takes  $O(\underline{n + v \log k})$ .

# Hash Tables Problem

## The very first question in LeetCode - TwoSum

---

- Given an array of integers, return indices of the two numbers such that they add up to a specific target.
- You may assume that each input would have exactly one solution, and you may not use the same element twice.
- Example: Given `nums = [2, 11, 7, 15]` and `target = 9`. Because `nums[0] + nums[2] = 2 + 7 = 9`, return `[0, 2]`.

## 1. Two Sum

2

11

7

15



# Hash Table in STL

Map, multimap, unordered\_map, HashMap

---

- Useful functions of STL map, multimap, unordered\_map
  - `size()`, `begin()`, `end()`, `empty()`
  - `insert(keyvalue, mapvalue)`
  - `find()`
  - `operator[]`
- Internal implementation:
  - map/multimap: Red-black tree
  - **unordered\_map**: Hash table

```
class Solution {
public:
    vector<int> twoSum(vector<int>& nums, int target) {
        unordered_map<int, int> m;
        for(int i = 0; i < nums.size(); i++)
        {
            if(m.find(target-nums[i]) != m.end())
                return {m[target-nums[i]], i};
            m[nums[i]] = i;
        }
    }
};
```

- Exercise problems from **Worksheet 8** (see “LA worksheet” tab in CS32 website). Answers will be posted next week.
- Questions for today: 1, 4, 2 (simplified 2Sum), 8 (3sum)

# Question 1

Given an array of distinct elements and a range [low, high], use a hash table to output all numbers in the range that are not in the array. Print the missing elements in sorted order. (F.Y.)

Example:

Input: `arr[] = {10, 12, 11, 15}, low = 10, high = 15`

Output: `13, 14`

Input: `arr[] = {1, 14, 11, 51, 15}, low = 50, high = 55`

Output: `50, 52, 53, 54`

## Question 2

Given a string, find the first non-repeating character in it and return its index. If it doesn't exist, return -1. You may assume the string contains only lowercase letters. Use a hash table to solve this problem.

Examples:

`s = "leetcode"`

`return 0`

`s = "loveleetcode",`

`return 2`

## Question 3

Given an array of integers and a target sum, determine if there exists two integers in the array that can be added together to equal the sum.

Examples:

Input:            arr[] = [4, 8, 3, 7, 9, 2, 5], target = 15

Output:           true

Explanation:    8 and 7 add up to the target sum 15

Input:            arr[] = [1, 3, 5, 2, 4], target = 10

Output:           false

Explanation:    No combination of two numbers in the array sum to 10

## Question 4

Write a function, `sum3`, that takes in an array of integers and determines whether there exists exactly three elements in the array that sum to 0. Return `true` if three such elements exist and `false` if not. No repeated elements are allowed. Your function must run faster than the brute force  $O(N^3)$ .

i.e `[1,2,3,4,5,6]` -> `False`

`[1,-1,2,-2]` -> `False`

`[1,2,-3, 6, 8]` -> `True`